

Wenhao Zhan

[Homepage](#) [Google Scholar](#)

Email: wenhao.zhan@princeton.edu

EDUCATION

Princeton University

Aug 2021-2026 (Expected)

Ph.D. Student in Electrical and Computer Engineering

Advisors: Jason D. Lee, Yuxin Chen

Princeton University

Aug 2021-Aug 2023

Master of Arts in Electrical and Computer Engineering

Advisors: Jason D. Lee, Yuxin Chen

- **Academic:** Overall GPA 4.00/4.00.

Tsinghua University

Bachelor of Electronic Engineering

Aug 2017-Jul 2021

- **Academic:** Major GPA 3.97/4.00, Overall GPA 3.92/4.00, rank 1/242.

RESEARCH INTERESTS

Reinforcement Learning, Large Language Models, Statistics, Optimization.

PUBLICATIONS&PREPRINTS

* = equal contributions, + = equal contributions and ordered randomly, # = equal contributions and ordered alphabetically

- K. Brantley, M. Chen[#], Z. Gao[#], J. D. Lee, W. Sun, **W. Zhan[#]**, X. Zhang, "Accelerating RL for LLM Reasoning with Optimal Advantage Regression ", 2025. Preprint.
- **W. Zhan**, S. Fujimoto, Z. Zhu, J. D. Lee, D. R. Jiang, Y. Efroni, "Exploiting Structure in Offline Multi-Agent RL: The Benefits of Low Interaction Rank ", 2025. The 13th International Conference on Learning Representations.
- A. Huang, **W. Zhan**, T. Xie, J. D. Lee, W. Sun, A. Krishnamurthy, D. J. Foster, "Correcting the Mythos of KL-Regularization: Direct Alignment without Overoptimization via Chi-squared Preference Optimization", 2025. The 13th International Conference on Learning Representations, **Spotlight**.
- Z. Gao, **W. Zhan**, J. D. Chang, G. Swamy, K. Brantley, J. D. Lee, W. Sun "Regressing the Relative Future: Efficient Policy Optimization for Multi-turn RLHF", 2025. The 13th International Conference on Learning Representations.
- J. D. Chang^{*}, **W. Zhan^{*}**, O. Oertell, K. Brantley, D. Misra, J. D. Lee, W. Sun, "Dataset Reset Policy Optimization for RLHF", 2024. Preprint.
- Z. Gao, J. D. Chang, **W. Zhan**, O. Oertell, G. Swamy, K. Brantley, T. Joachims, J. A. Bagnell, J. D. Lee, W. Sun, "REBEL: Reinforcement Learning via Regressing Relative Rewards", 2024. The 38th Conference on Neural Information Processing Systems.
- Z. Zhang, **W. Zhan**, Y. Chen, S. S. Du, J. D. Lee, "Optimal Multi-Distribution Learning", 2024. The 37th Annual Conference on Learning Theory.
- **W. Zhan**, M. Uehara, W. Sun, J. D. Lee, "Provable Reward-Agnostic Preference-Based Reinforcement Learning", 2024. The 12th International Conference on Learning Representations, **Spotlight**.

Wenhao Zhan

[Homepage](#) [Google Scholar](#)

Email: wenhao.zhan@princeton.edu

- **W. Zhan***, M. Uehara*, N. Kallus, J. D. Lee, W. Sun, "Provable Offline Preference-Based Reinforcement Learning", 2024. The 12th International Conference on Learning Representations, **Spotlight**.
- Y. Zhao⁺, **W. Zhan**⁺, X. Hu⁺, H. Leung, F. Farnia, W. Sun, J. D. Lee, "Provably Efficient CVaR RL in Low-rank MDPs", 2024. The 12th International Conference on Learning Representations.
- G. Li*, **W. Zhan***, J. D. Lee, Y. Chi, Y. Chen, "Reward-agnostic Fine-tuning: Provable Statistical Benefits of Hybrid Reinforcement Learning", 2023. The 37th Conference on Neural Information Processing Systems.
- **W. Zhan***, S. Cen*, B. Huang, Y. Chen, J. D. Lee, Y. Chi, "Policy Mirror Descent for Regularized Reinforcement Learning: A Generalized Framework with Linear Convergence", 2023. SIAM Journal on Optimization.
- **W. Zhan**, M. Uehara, W. Sun, J. D. Lee, "PAC Reinforcement Learning for Predictive State Representations", 2023. The 11th International Conference on Learning Representations.
- **W. Zhan**, J. D. Lee, Z. Yang, "Decentralized Optimistic Hyperpolicy Mirror Descent: Provably No-Regret Learning in Markov Games", 2023. The 11th International Conference on Learning Representations.
- **W. Zhan**, B. Huang, A. Huang, N. Jiang, J. D. Lee, "Offline Reinforcement Learning with Realizability and Single-policy Concentrability", 2022. The 35th Annual Conference on Learning Theory.
- C. Z. Lee, L. P. Barnes, **W. Zhan**, A. Özgür, "Over-the-Air Statistical Estimation of Sparse Models", 2021. The 2021 IEEE Global Communications Conference.
- **W. Zhan**, H. Tang, J. Wang, "Delay Optimal Cross-Layer Scheduling Over Markov Channels with Power Constraint", 2020. The IEEE International Symposium on Broadband Multimedia Systems and Broadcasting 2020.

REVIEWING EXPERIENCES

- ICML 2022, 2023, 2024
- NeurIPS 2022, 2023, 2024
- ICLR 2024
- OPT 2022
- Operations Research (INFORMS)
- Mathematical Programming (Springer)
- Machine Learning (Springer)
- SIAM Journal on Mathematics of Data Science

WORKING EXPERIENCES

Meta GenAI

Jun 2025-Present, Research Intern

Project: Reinforcement Learning For Tool-Integrated Reasoning Models

- Explored the ability of large language models in adaptively leveraging external tools in long-context reasoning.
- Developed a novel **three-stage post-training pipeline** for mid-trained models, incorporating few-shot fine-tuning, tool-usage-oriented reinforcement learning, and answer-oriented reinforcement learning.
- Achieved over 10% improvement in performance and set new state-of-the-art results on several challenging mathematical reasoning benchmarks.

Wenhao Zhan

[Homepage](#) [Google Scholar](#)

Email: wenhao.zhan@princeton.edu

Meta Ranking

May 2024-Oct 2024, Research Intern

Project: Efficient Multi-Agent Offline Reinforcement Learning

- Investigated the coordination among multiple agents in complex environments, such as auctions and robotics.
- Identified a key structural property of the network architecture—**interaction rank**—and established a theoretical framework for decentralized, computationally and statistically efficient offline reinforcement learning.
- Conducted experiments demonstrating the effectiveness of critic architectures with low interaction rank, showing substantial performance gains over standard single-agent value decomposition baselines used in prior work.

TEACHING EXPERIENCES

- **Foundations of Reinforcement Learning**
TA, Spring 2024, Princeton University
- **Special Topics in Information Sciences and Systems: Theory of Deep Weakly Supervised Learning**
TA, Fall 2022, Princeton University

HONORS & AWARDS

- 2024 Award for Excellence by Princeton SEAS
- Honorable mention for the 2023 Jane Street Graduate Research Fellowship
- 2017-2020 Tsinghua Academic Excellence Award
- 2018-2020 Tsinghua Scientific Research Excellence Award
- 2018-2020 National Encouragement Scholarship

INVITED TALKS

- **Optimal Multi-Distribution Learning**
Adaptive Learning in Complex Environments, TTIC Chicago Summer Workshop 2024

TECHNICAL SKILLS

- Programming languages: C/C++, Python, Matlab, Verilog
- Tools: Git, PyTorch
- Cloud: AWS, SLURM